

CNN AND TRANSFER LEARNING

ASSIGNMENT 3

CNN

1. Hyperparameters

- a. Learning rate
 - i. 0.01, 0.001
- b. Optimizer
 - i. sgd, adagrad, adam, rmsprop
- c. CNN layers
 - i. 1, 2, 3
- d. Pooling layers
 - i. 1, 2, 3
- e. Pooling type
 - i. AveragePooling2D, MaxPooling2D, GlobalMaxPooling2D, GlobalAveragePooling2D
- f. Regularization
 - i. Early stopping, l1, l2, data augmentation

2. Optimizer and LR HP (using best model found in step 3)

Optimizer	SGD		Adagrad		RMSprop		Adam	
LR	0.01	0.001	0.01	0.001	0.01	0.001	0.01	0.001
Train-Acc	0.97	0.96	0.98	0.98	0.95	0.97	0.97	0.98
Test-Acc	0.91	0.93	0.95	0.94	0.90	0.94	0.90	0.95

3. CNN and pooling layers HP

CNN → 7x7 filters with 1x1 stride length and same padding for first layer

MaxPool → 2x2 with 1x1 stride length

Using Batch normalizer for (2 CNN, 1 MP) (2 CNN, 2 MP) (3 CNN, 3 MP)

CNN Layers	1 CNN		2 CNN		3 CNN	
Pooling Layers	1	2	1	2	2	3
Train-Acc	0.94	0.92	0.99	0.98	0.91	0.97
Test-Acc	0.91	0.86	0.953	0.93	0.90	0.95
Model	m_1	m_2	m_3	m_4	m_5	m_6

4. Pooling type HP

Pooling Type	Avg	Max	Glob_Avg	Glob_Max
Train-Acc	0.96	0.98	-	-
Test-Acc	0.91	0.953	-	-

5. Regularization HP

Reg	None	ES	L1	L2
Train-Acc	0.98	0.96	0.95	0.97
Test-Acc	0.953	0.94	0.91	0.92

6. Observation

- Best Model

-
- i. Optimizer Adam
 - ii. LR 0.001
 - iii. No regularization
 - iv. Use Batch normalization
 - v. 2 CNN layers and 1 MaxPool layer
- b. Training accuracy is always higher than test accuracy
 - c. Test accuracy is varying due to small epoch
 - d. Best Accuracy is 0.95

ResNet & VGG

1. Untrained

a. Hyperparameters

- i. Optimizer
 - 1. Adam, SGD
- ii. LR
 - 1. 0.01, 0.001
- iii. Regularization
 - 1. Early stopping, l1, l2 (choose random Conv2D layers to apply kernel regularization)
- iv. Pooling
 - 1. Average, max, none
- v. Epochs
 - 1. 10, 20, 40
- vi. Top-layer
 - 1. Adding dense (512, 1024) with dropout
 - 2. Adding dense (128, 512, 1024) without dropout

b. Hyperparameter tuning

(Using **SGD**(0.0001, 0.9 mom) and dense layer 128 neuron))

ResNet

Epochs	10		20		40	
Pooling	max	avg	max	avg	max	avg
Train-Acc	0.85	0.85	0.90	0.91	0.87	0.96
Test-Acc	0.64	0.56	0.86	0.9	0.91	0.9

(Using **Adam**(0.001) and dense layer 128 neuron))

ResNet												
Epochs	10				20				40			
Pooling	max		avg		max		avg		max		avg	
LR	0.001	0.01	0.001	0.01	0.001	0.01	0.001	0.01	0.001	0.01	0.001	0.01
Train-Acc	0.96	0.59	0.98	0.51	0.82	0.52	0.95	0.52	0.96	0.52	0.96	0.52
Test-Acc	0.46	0.58	0.56	0.54	0.49	0.56	0.94	0.56	0.90	0.56	0.87	0.56

(Using Avg pooling and 20 epochs)

ResNet												
Dnese	128				512				1024			
Dropout	No		0.5		No		0.5		No		0.5	
Opt	Ada m	SGD	Ada m	SGD	Ada m	SGD	Ada m	SGD	Ada m	SGD	Ada m	SGD
Train-Acc	0.93	0.89	0.72	0.86	0.88	0.93	0.95	0.89	0.95	0.91	0.97	0.87
Test-Acc	0.80	0.82	0.88	0.86	0.84	0.84	0.93	0.85	0.89	0.82	0.88	0.84

Model	ResNet152
Optimizer	Adam
Dense	512
LR	0.001
Pooling	Average

Epochs	20			
Dropout	0.5			
Reg	None	ES	l1	l2
Train-Acc	0.97	0.91	0.86	0.96
Test-ACC	0.91	0.54	0.57	0.92
Training Time	79 sec	58 sec	79 sec	79 sec
TT per epoch	3 sec	3 sec	3 sec	3 sec
Test time	0.34 sec	0.35 sec	0.37 sec	0.34 sec
Multiplication	Takes too much time to compute			

Model	VGG16							
Optimizer	Adam				SGD			
LR	0.001		0.0001		0.0001		0.00001	
Reg	L2	Drop	L2	Drop	L2	Drop	L2	Drop
Train-Acc	0.72	0.78	0.96	0.97	0.80	0.66	0.85	0.62
Test-ACC	0.56	0.64	0.68	0.96	0.71	0.67	0.82	0.70
Training Time	23 sec							
TT per epoch	1 sec							
Test time	0.1 sec	0.1 sec	0.12 sec	0.19 sec	0.16 sec	0.11 sec	0.1 sec	0.1 sec
Multiplication	$(64*64*64*3*3*3)+(32*32*128*3*3*64)+(16*16*256*3*3*128)+(8*8*512*3*3*256)+(4*4*512*3*3*512)+(2*2*512*3*3*512)+(2048*4096)+(4096*4096)+4096$							

2. Pretrained HP tuning

-
- a. Number of freezed layers
 - b. Epochs
 - c. Number of dense layers (1 or 2)
 - d. Neurons in dense layer (512, 512 and 256)
 - e. Optimizers (RMSprop, Adam, SGD)
 - f. Learning rate (0.001, 0.0001)
 - g. Regularizers (L1, L2, Dropout)

3. Best result achieved for:

- a. ResNet : (opt-Adam, LR-0.001, Reg-L2/no reg, dense layers 2 (512, 256), epochs 35, no freezed) most stable highest accuracy -> 0.94/0.86 followed by 20 layer freezing with 0.90 (by trying 5,15, 20,30,60,100)

There is excessive randomness during runs I don't know why, same combinations gave different outputs when run repeatedly

[Layer -Freezing plot](#)

[Accuracy Without freezing](#)

- b. VGG: (opt-Adam, LR-0.0001, no reg, dense layers 2 (512, 256), epochs 25, no freezed) most stable highest accuracy -> 0.973 followed by 5 layer freezing with 0.98 (by trying 1,2,5,10,15)

[Layers-Freezing plot](#)


[Accuracy Without freezing](#)

[Accuracy with freezing](#)

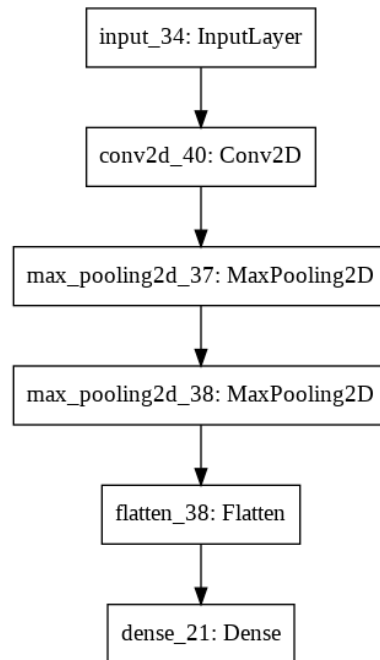
Appendix

1. CNN

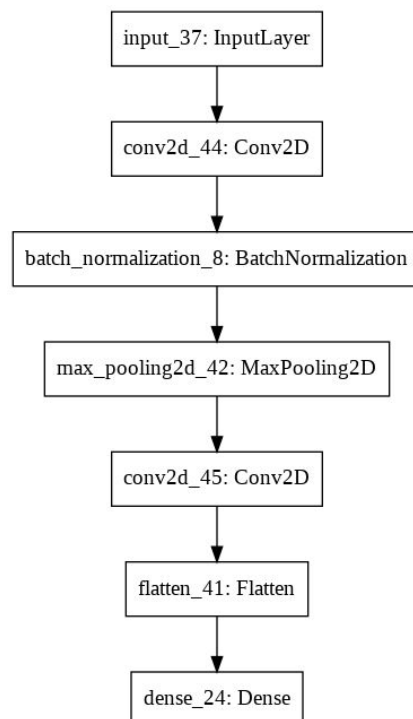
- a. LR and Optimizer tuning
- b. Model CNN and MP layers tuning
 - i. M_1



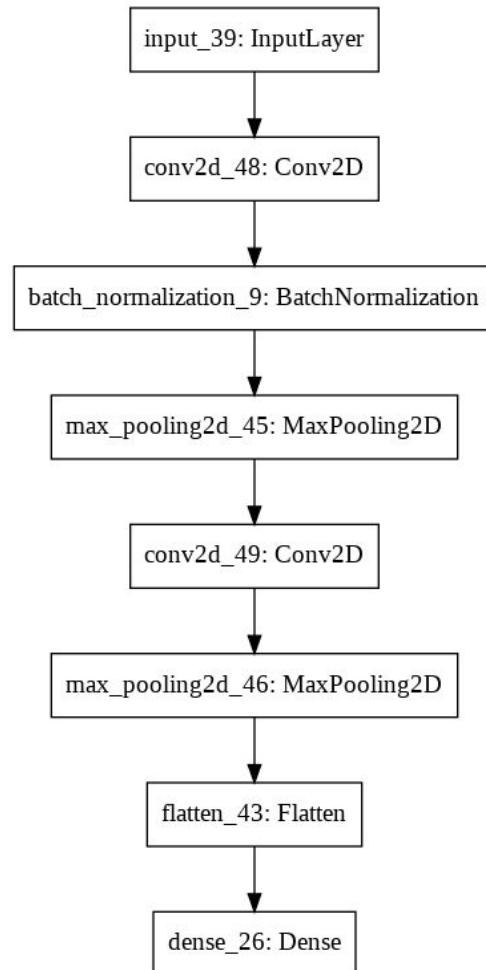
ii. M_2



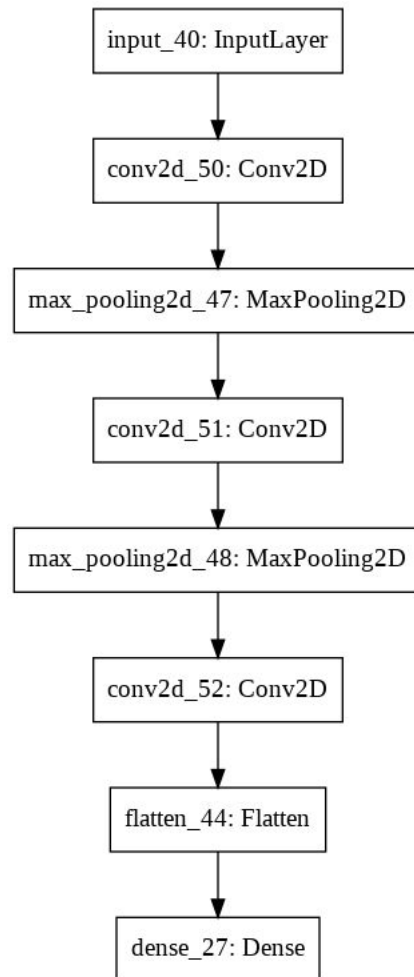
iii. M_3

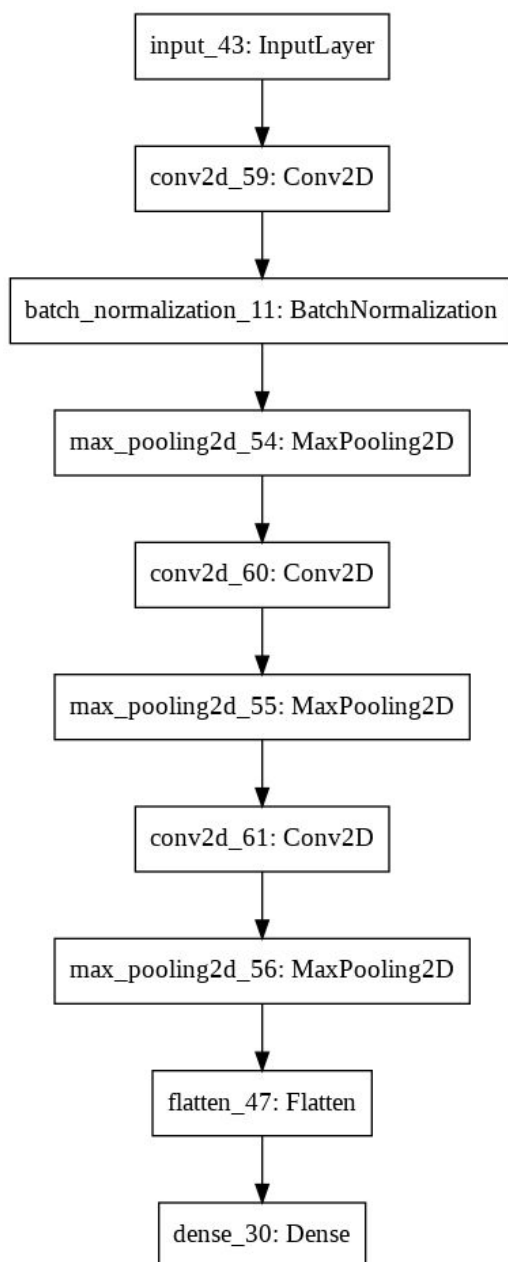


iv. M_4




v. M_5





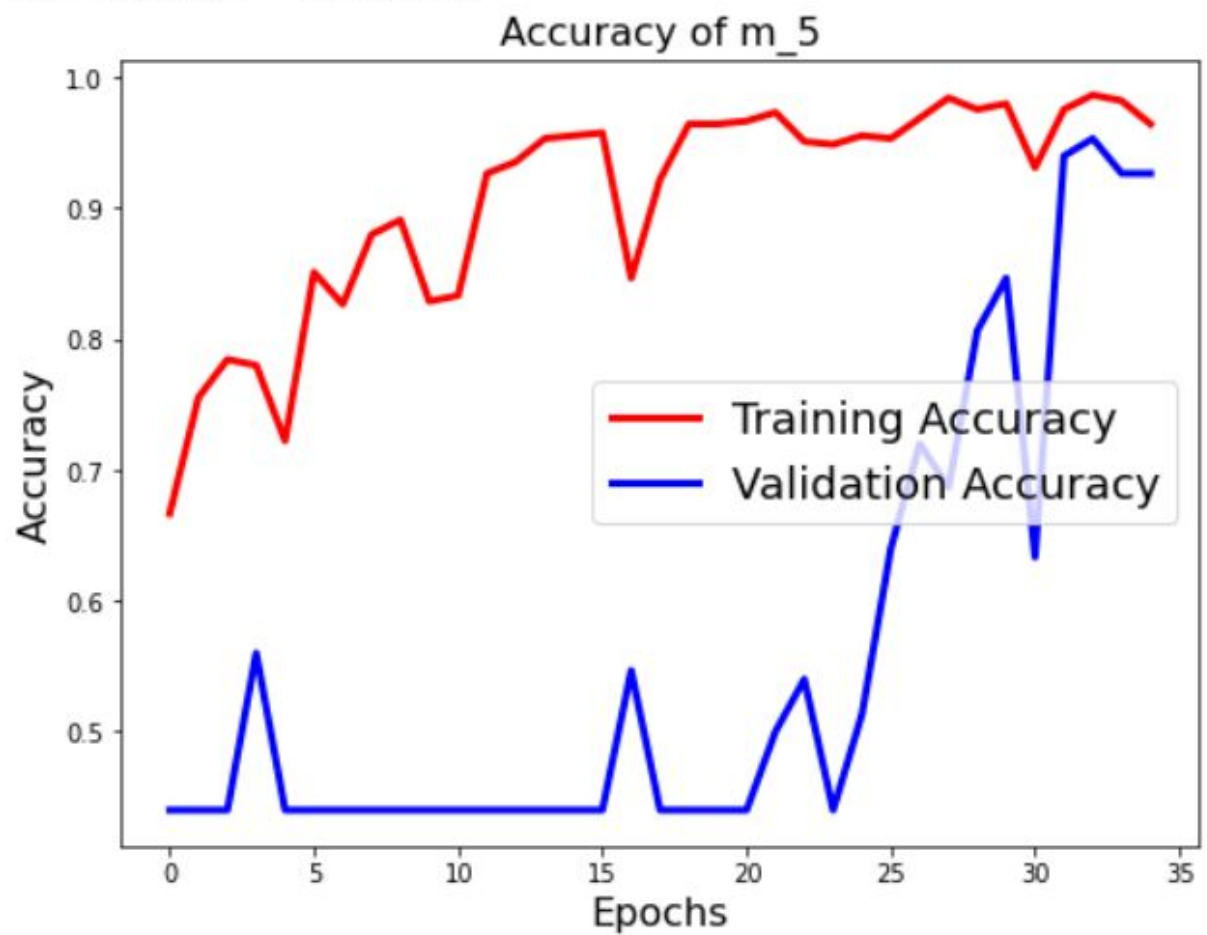
c. Freezing layer plot ResNet



d. Without freezing Resnet

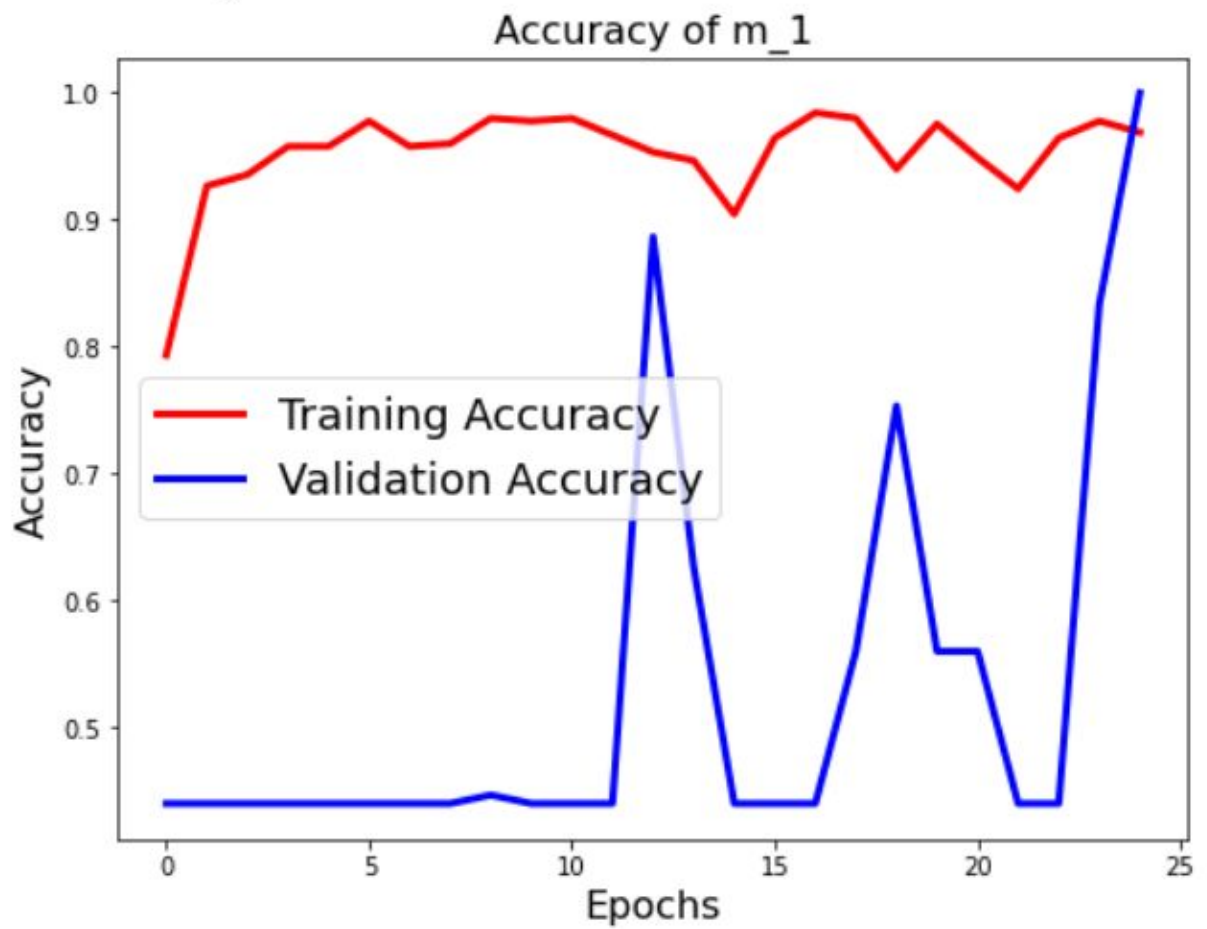
Loss = 0.15734563946723937

Test Accuracy = 0.94666666



e. Without freezing VGG


Test Accuracy = 0.97333336



f. With freezing VGG

Loss = 0.07534479891260465

Test Accuracy = 0.98



g. Freezing layer plot VGG

